

TradeLens

Post-trade Review · Problem 03

Name Shraman Hazra
Email hazrashraman2002@gmail.com
Phone 8900406041

LIVE PROTOTYPE

<https://app-azure-zeta-42.vercel.app>

SOURCE CODE

<https://github.com/Shraman123/tradelens>

CONTENTS

1. The one-pager
2. The prompts
3. The evals
4. The working app

ARTIFACT 1 OF 4

The one-pager

TradeLens — a post-trade review that names the habit, not the metric

Problem 03 • Post-trade review · Live prototype: <https://app-azure-zeta-42.vercel.app>

The problem

My user is an intraday index-options buyer on Nubra. He trades most days and reviews almost never — not a data problem, since his order history already sits inside the app, but a diagnosis problem: after a losing month he knows the number, not which behaviour produced it.

SEBI's August 2026 studies: 87.7% of individual F&O traders lost money in FY26, and about 90% of those who lost two years running and kept trading lost again the next year. Losing is not the anomaly. Not changing is.

So TradeLens answers not “how did I trade last month” but “what am I repeatedly doing that costs me money”.

Why Nubra should solve it

Nubra already holds the order book — no upload, no journalling discipline, no delay; the review can arrive the morning after a bad week, not whenever the trader gets round to it.

A broker earns on volume, and this product will often tell a trader to trade less. I think it is right to build: a churned trader generates no volume, and diagnosing one costly habit protects the account that produces future revenue.

The solution space

Four rejected shapes: a dashboard shows metrics and leaves diagnosis to the user; a journal demands discipline he has already shown he lacks; a chatbot answers questions he does not know to ask; real-time nudges need infrastructure and trust this product has not yet earned.

What I built: a monthly review naming at most three habits, ranked by cost, each proved with the trader's own trades, ending in one process rule he commits to. Next month's review opens with how often he broke it. The review is the artifact; the loop is the product.

How it works

Code computes every number; the model only explains the JSON it produces and may not state a figure not in it. A habit clears three gates: enough trades, statistical significance against the trader's own other trades, and cost above a materiality bar tied to his turnover. Every cost is a counterfactual, not a raw loss. Output has four states: confirmed habit, watching, nothing found, not enough data. Advice is process only — size, timing, trade count, stops — never instruments or direction.

What testing changed

I wrote the evals before the app, using synthetic traders with habits planted at known strength, plus a control and a thin-data trader. One month of history recovered the planted habits in only 13 and 20 of 30 runs; 60 days recovered them in 28 and 29 of 30, with zero false alarms on the control. Loosening the

significance threshold bought recall only by inventing habits — so the window became 60 days, and unproven-but-material patterns are shown as “watching,” not overstated.

What I’d validate first

This rests on assumptions untested with a real trader. Three matter most.

Desirability: a losing trader wants diagnosis, not avoidance. Test: do reviews open after losing months, or only winning ones? If opens track wins, the review needs pushing, not waiting to be opened.

Behaviour change: naming a habit and a commitment reduces it next month. Test: recurrence, committed vs not. If it’s the same either way, the loop does not work.

Trust: a trader accepts a verdict from a broker who earns on his volume. Test: commit rates with and without a plain disclosure of that incentive. A gap means real distrust, and a different messenger.

How I’d measure success

The metric is habit recurrence next month, not opens or session time. The counter-metric is churn — if diagnosed traders leave rather than improve, the thesis is wrong. A trader can also fix the named habit and still have a worse month; one habit is rarely the only cause. Success is measured on the habit, not the P&L.

Scope

In: five screens, seven detectors, four output states, synthetic order history for four traders, and an LLM narration layer with automated evals for number tracing, advice, state fidelity and stability.

Out: live market data, real user data, price behaviour after exit, real-time nudges, instrument or direction advice, and price prediction.

ARTIFACT 2 OF 4

The prompts

Prompt log — summary

Full log, including every failed narration-prompt version, follows this page. Numbered sections below match PROMPTS_LOG.md's own numbering.

#	Date	What it covers	What changed
1	2026-09-11	Initial build brief (Claude Code)	Product scoped: 5 screens, 4 personas, code-computes/LLM-narrates architecture, no-advice hard line.
2	2026-09-11	Go-ahead + provider switch	Approved tool installs; narration LLM switched to Groq (free tier) after the Anthropic key had no credit.
3	2026-09-11	Narration prompt v1–v5	v1 FAILED (raw floats, not plain English). v2 FAILED (self-modelled the banned word “habit”). v3 FAILED (closing nudged toward a rule even with none to pick). v4 passed automated checks. v5 fixed why_it_matters completeness for rerun stability.
4	2026-09-11	Reliability incident, v5 stability run	Malformed-JSON response and a digit-insertion flake found mid-batch; added a runtime self-check retry and reliability counters.
5	2026-09-12	Four fixes from the deployed app	Not-reported list humanized (labels + plain-English reasons); Vikram's closing rewritten; window wording fixed; habit-detail cost relabeled; concrete rule for size_up_after_loss.
6	2026-09-12	Narration prompt v6	“This window” not “this month”; one narrow exception letting the nothing-found closing state the trade count; concrete-rule builder. Stability: 5/5, all four personas.
7	2026-09-12	Two copy fixes, v7	Ruled-out phrasing varied by p-value tier instead of one repeated sentence; fixed a false-causal “so” in the nothing-found closing. Stability: 5/5.
8	2026-09-12	Live-app bug: “a loss of a loss”, v8	A duplicated phrase automated checks couldn't catch (correct number, broken English) found by reading the live app; new prompt rule plus a permanent redundant_framing check.
9	2026-09-12	v8 stability, partial	Arjun 5/5, Neha 1/5, Vikram/Sara blocked by Groq's daily quota. Also fixed eval_narration.py silently reporting a missing stability run as a pass instead of a skip.
10	2026-09-12/13	ONE_PAGER.md + PDF pipeline	Added the one-pager (699 words) and a pandoc + Puppeteer build for this submission PDF.

Full prompt log

TradeLens — Prompt Log

Every prompt given to Claude Code in this repo, in order, including ones that didn't work. Prompts to the narration LLM (in prompts/) are logged here too as they're written.

1. Initial build brief (Claude Code, 2026-09-11)

```
# Claude Code prompt – build the TradeLens prototype
```

```
Unzip `tradelens_detectors_and_evals.zip`, open Claude Code inside the `tradelens/` folder, and paste everything below the line as your first message. Paste it into your prompt log too; the assignment asks for every prompt you used.
```

```
---
```

```
You are building TradeLens, a clickable prototype for a product-management take-home. Read this whole brief, then read `habit_library.md`, `eval_report.md`, `detectors.py` and everything in `demo/` before writing any code. Then give me a short plan and wait for my go-ahead.
```

```
## What the product is
```

```
A post-trade review for intraday index-options buyers on Nubra, an Indian broker. It turns 60 days of order history into the two or three habits that are costing the trader money, proves each one with the trader's own trades, prices the habit with a counterfactual, and gets the trader to commit to one process rule for next month.
```

```
Product principle: insight over information. This is not a dashboard, not a journal, not a chatbot. If a screen shows a metric that doesn't change what the trader should do next, cut it.
```

```
Core architecture, which is non-negotiable:
```

```
- Code does the math. `detectors.py` has already been run; its outputs are the `demo/*_analysis.json` files. The app never recomputes or invents a number.
```

```
- The LLM does the words. It turns one analysis JSON into plain-language narration. It may not state any number, count or percentage that is not present in the JSON.
```

```
- Process advice only. Rules are about size, timing, trade count and stops. Never suggest an instrument, strike, direction or "buy/sell" anything. Nubra is a SEBI-registered broker; this is a hard line.
```

```
## The four output states (all must be clickable)
```

```
`demo/` has four personas, and each shows a different state. Do not pick flattering data; use the files as they are.
```

```
| Persona | State | What the screen must do |
```

```
|---|---|---|
```

```
| `arjun_revenge_sizer` | Confirmed habit + one "watching" pattern | Show habit #1 with evidence and cost; show averaging-down as watching, clearly not a habit |
```

```
| `neha_expiry_day` | One confirmed habit | Show it; the "not reported" list explains what was ruled out and why |
```

```
| `vikram_control` | Nothing found | Say plainly he's losing money but no repeated pattern was found; no fake habit, no filler |
```

```
| `sara_thin_data` | Not enough data | Say how many trades she has and the minimum; no habits |
```

```
The persona switcher is the demo device: a reviewer clicks from Arjun to Vikram and watches the product refuse to
```

invent a problem.

Screens

1. **Persona switcher** (top of the app, always visible). Name, one-line description, and what state they demonstrate.
2. **Review**. Review month (August 2026) in three numbers: net P&L, trades, win rate. Evidence window stated explicitly ("based on 1 Jul-31 Aug"). Then the habits, ranked, each as a card: label, cost, confidence, one-sentence "why it matters", one process rule. Then "Watching" (if any), then "Also noticed" (uncosted flags), then a collapsed "What we checked and ruled out" list from ``not_reported``.
3. **Habit detail**. The evidence table from ``evidence``, the counterfactual sentence, the five ``example_trades`` with a link to view each row order in the CSV, and the rule. State that habit costs overlap and are never added together.
4. **Commit**. The trader picks one rule (the top habit's is pre-selected, others available) and confirms it. One rule only.
5. **Next month preview**. A mocked version of next month's review opening with "You broke this rule 3 times in 21 days" so the reviewer sees the loop, not just the report. Label it clearly as a preview of the loop, not real data.

Empty and thin-data states are screens, not error toasts. Design them as carefully as the happy path.

Narration layer

Write ``narrate.py`` and ``prompts/narration_system.md``. The script sends each analysis JSON to Claude and writes ``demo/<persona>narration.json`` with: a headline, one paragraph per habit ("why it matters"), one rule per habit, and a short closing. Commit the generated narrations so the deployed app is static and fully clickable with no API key. Keep a "regenerate" command in the README.

The system prompt must enforce:

- every number in the output appears verbatim (or as the same value rounded) in the input JSON;
- no instrument, strike, direction or market prediction;
- "watching" items are described as something to keep an eye on, never as a habit;
- "not enough data" and "nothing found" are stated honestly, with no consolation habits invented;
- plain English, no jargon a first-year trader wouldn't know, no moralising.

Then write ``eval_narration.py`` that checks all of that automatically:

- **Number tracing**: extract every number from the narration and assert it exists in the JSON. Fail on any orphan number.
- **Advice filter**: fail on any instrument/strike/direction words (build a word list: CE, PE, call, put, buy, sell, short, long, strike, Nifty, Bank Nifty, expect, will rise, will fall, target).
- **State fidelity**: for Vikram and Sara the narration must contain no habit language; for Anjun the watching item must not use "habit"/"pattern you have".
- **Stability**: run narration 5 times per persona and assert the set of habits named and every number are identical each run.

Run it and report results. If any check fails, fix the prompt, not the checker, and keep the failed prompt version in ``prompts/`` with a note on why it failed. I need a full record of every prompt that didn't work.

Stack and deployment

- Single-page React app with Vite, TypeScript, Tailwind. No backend; load the JSON files statically.
- Deploy to Vercel (or equivalent free static host) with a public URL. The assignment requires a link a reviewer can open without login.
- Mobile-first: most Nubra users are on a phone. Must also look right on desktop.
- Dark UI, restrained, feels like a broker's app and not a college project. One accent colour for "cost". Currency formatted as ₹1,05,732 (Indian grouping). No emojis in the UI.
- No charts unless a chart is the evidence itself. A comparison of two numbers is two numbers, not a bar chart.

How I want you to work

- Plan first, then build screen by screen. After each screen, stop and show me what changed.
- Do not change anything in ``detectors.py``, ``generate_personas.py`` or the JSON files. If you think a number or label in the data is wrong, tell me; don't patch it.
- Keep ``PROMPTS_LOG.md`` at the repo root and append every prompt I give you and every prompt you write for the LLM, including ones that didn't work.
- Write a short ``README.md``: what it is, how to run, how to regenerate narrations, how to run both eval scripts, and a list of every assumption (lot size, expiry day, charges model, synthetic data).
- Ask me before adding any feature not listed here. The bar is "clickable and honest", not "complete".

Start by reading the files and giving me the plan.

2. Go-ahead prompts (Claude Code)

- yes — approved installing Node/Git/Python via winget and starting on the narration layer.
- On being asked how to handle the missing `ANTHROPIC_API_KEY`, chose “you provide a key” (real API calls, not self-authored narration) over the fallback option.
- no billing bro cab i use free ai keys any groq or deepseek — after the Anthropic key came back with insufficient credit, asked to switch to a free provider. Result: added a `--provider` switch to `narrate.py` (`anthropic / groq`), defaulting to Groq (genuinely free, no card; DeepSeek is prepaid like Anthropic so was ruled out). See `narrate.py`’s `call_groq`.

3. Narration system prompt iterations (the LLM-facing prompt, not a Claude Code prompt)

Each version is saved in full under `prompts/`. Summary of what changed and why (the “why” for each failure has its own `*_FAILED.notes.md` next to the saved version):

Version	File	Result
v1	<code>narration_system_v1_FAILED.md</code>	Failed (caught by manual read, not by <code>eval_narration.py</code> ’s 4 checks): sent the model raw evidence floats and told it not to invent numbers, but never said how to format the numbers it was allowed to use. It pasted JSON floats verbatim (“cost you 105732.0”, “win rate of 0.281”) — traceable, but not plain English, and not something a broker app should show a user. Also: Sara’s (insufficient-data) closing used the word “habit” while honestly explaining there wasn’t one yet, failing the state-fidelity check as specified.
v2	<code>narration_system_v2_FAILED.md</code>	Failed (state_fidelity, Arjun): fixed v1’s two failures by having <code>narrate.py</code> pre-format every

Version	File	Result
		number into its exact display string (₹ Indian grouping, %, ratios) and telling the model those are text to copy, not numbers to interpret. But rule 3's own example text said "not yet a confirmed habit" — the model echoed that phrase almost verbatim, word "habit" included. Self-inflicted: I demonstrated the exact word I was telling it to avoid.
v3	narration_system_v3_FAILED.md	Failed (caught by manual read, not automated): fixed v2 by explicitly banning the literal word "habit" in watching notes, even negated. All 4 automated checks passed for all 4 personas — but Vikram's (nothing-found) closing read "Pick one rule to focus on for next month," and Vikram has zero qualified habits. The closing instruction was unconditional ("nudging toward picking one rule") when it should only apply when a habit exists to pick. Passed. Closing instruction made conditional on habits being non-empty; when empty, forbids "pick/choose/focus on/apply a rule" language and asks for an honest, encouraging line instead. All 4 personas pass all 4 eval_narration.py checks (number tracing, advice filter, state fidelity; stability run separately, see eval_report.md).
v4	narration_system.md (current)	

v5 | narration_system.md (current) | **Passed** number tracing / advice filter / state fidelity for all 4 personas. First 5x stability run (real python narrate.py --stability 5, 20 Groq calls) still in progress at time of writing — see eval_report.md for the final numbers once complete. |

v4's 5x stability run (also real, 20 Groq calls) found genuine content-selection variance, not random noise: why_it_matters cited a different subset of a habit's evidence fields each run (nothing invented, just inconsistent which facts appeared), a watching note sometimes stated a cost figure and sometimes didn't, and Sara's insufficient-data headline sometimes restated the trade count and sometimes didn't — v4 left all three optional. v5's fix: for a habit, why_it_matters must cite the cost *and every* evidence field, every time (completeness, not curation); watching/also_noticed notes must never contain a number, full stop; headline/closing must never contain a number, full stop (the UI shows exact figures directly). Removing

the optionality is what removing the instability required — an LLM asked to freely choose what to include will choose differently run to run by construction. Full detail in `prompts/narration_system_v4_FAILED.notes.md`.

Provider/model used throughout: Groq, openai/gpt-oss-120b, temperature 0.2.

4. Reliability incident during the v5 stability run (2026-09-11)

Mid-batch (run 5/5 for `neha_expiry_day`, ~9 calls into a 20-call stability run), Groq returned a 200 OK whose message content wasn't valid JSON. `narrate.py` only retried HTTP-level failures at the time, not a malformed-*content* response inside a successful HTTP call, so it crashed the whole batch instead of retrying that one call.

Fix, two parts: 1. Moved JSON parsing inside the retry loop in `_post` (a bad-content response now retries like any other transient failure), and separately added a self-check pass in `narrate_persona` — after generating, validate the output against the same `check_number_tracing / check_advice_filter / check_state_fidelity` functions `eval_narration.py` uses, and regenerate (up to 2 extra attempts) if any fail. This is also what caught and would auto-retry the digit-insertion bug below. 2. Added STATS counters in `narrate.py` (top-level calls, malformed-JSON retries, network/HTTP retries, self-check retries) so future runs report a real malformed-output rate instead of a guess. **Caveat:** this specific incident predates the counter, so its exact rate isn't recorded — only that it happened once in roughly 60-70 calls across all runs so far (rough, not measured). The existing prompt already said “Return only a JSON object, no prose outside it” at the top of the Output Format section, so this reads as an infrequent model-reliability flake on the free tier, not a demonstrated prompt-content gap — I'm holding off on a v6 prompt reinforcement (“no markdown fences”, repeated JSON-only instruction) unless a future run's *measured* rate (via the new counters) comes in meaningfully above a few percent. Rerunning stability with the fix in place succeeded 20/20.

Separately, that same crashed run's completed portion surfaced a genuine content bug (not a crash): one of the 5 `neha_expiry_day` runs said “a loss of ₹9,449” for a field where every other run correctly said “₹449” — same `rest_avg` evidence value (-449.0), a digit inserted by the model despite the value being handed to it as a pre-formatted, copy-verbatim string. This is exactly what the new self-check retry (point 1 above) is meant to catch automatically going forward, and it's now covered in `eval_narration.py`'s own number-tracing logic being applied per-generation, not only at the end.

5. Four fixes from clicking through the deployed app (Claude Code, 2026-09-12)

Four fixes from clicking through the deployed app. Do them all, then build, lint, test, commit and push.

1. BUG — empty ruled-out list. On the Review screen, “What we checked and ruled out (7)” expands to 7 empty bullets (seen on Vikram, check all four personas). Each row should render the human-readable detector label plus its rejection reason from `not_reported` in the analysis JSON — e.g. “Sizes up after a loss — the size difference was within normal variation”. Map detector ids to the labels in `habit_library.md`, and turn the raw reason strings into plain English (a reader shouldn't see “`p=0.34 >= 0.01`” unless they expand further).

2. Vikram's closing line. “Keep trading as usual and check back next review” reads as “change nothing” to someone down ₹52,169. Replace with wording that says: nothing here is a habit yet; across 101 trades no pattern repeated often enough or cost enough to be worth changing; that doesn't mean the month went well, it means the loss isn't traced to one recurring behaviour we can prove from his orders. Keep “check back next review” as a separate, quieter line. If this text comes from the narration layer rather than the UI, change the prompt instead and

regenerate – don't hardcode it.

3. Window wording. Vikram's headline says "this month" but the window is 1 Jul–31 Aug. Make all headline/summary copy say "in this window" or "over these 60 days". Check every persona.

4. Habit detail – the cost number and the rule.

a. The large ₹ figure reads like a loss. Label it "What this cost you" above the number, with the counterfactual sentence underneath in smaller text.

b. Make the rule concrete using numbers already in the evidence. For `size_up_after_loss`: "After a loss, your next entry stays at your usual size – about ₹6,600. If you want to size up, it has to be a separate decision made before the session, not after a loss." Substitute the persona's own `median_entry_size_otherwise`, rounded. Apply the same treatment to any other detector whose evidence supplies a concrete threshold.

Constraints: don't touch `detectors.py`, `generate_personas.py` or the analysis JSONs. Any number appearing in new copy must already exist in the analysis JSON. After the changes, re-run `eval_narration.py` (number tracing, advice filter, state fidelity) and paste the results. Append every prompt you write for the LLM to `PROMPTS_LOG.md`, including versions that fail the evals.

Findings and decisions made while implementing this, before touching any file:

- **Fix 1 was code, not the LLM.** `not_reported` is rendered directly by `ReviewScreen.tsx` from `analysis.json` — it was never sent to the narration LLM (`narrate.py`'s `trimmed_input` deliberately excludes it, see its own comment). So "map ids to habit_library.md's labels" and "humanize the reason" are both a `habitLabels.ts` change, not a prompt change. The bullets were never literally empty in the deployed build (verified against the live JS bundle before this prompt) — the real complaint, confirmed by the brief's own example, is that a reader shouldn't have to decode $p=0.34 \geq 0.01$; that's what "empty" meant here. New: `humanizeReason()` in `habitLabels.ts`, pattern-matching every reason format `detectors.py` actually emits and rewording each into a plain sentence — never inventing a number, only dropping or rewording the ones already in the string. `HABIT_LABELS` itself now sources its wording from `habit_library.md` (previously it mirrored `detectors.py`'s own `label` field verbatim — a deliberate earlier choice, reversed here because this prompt explicitly asks for the other source). The "explained by X (confounded)" case looks up X's real, currently-displayed label from `analysis.habits` when possible, so the cross-reference never disagrees with the habit card it's pointing at.
- **Fix 3's "60 days" isn't literally true for every persona** — checked before writing the prompt: Arjun/Neha/Vikram's window is 62 days (2026-07-01 to 2026-08-31), Sara's is 59 (thin-data, truncated to her actual history). Hardcoding "60 days" into the prompt would be wrong for every one of the four, and would fail Sara's number-tracing (59 vs 60 is right at the ± 1 tolerance edge; 62 vs 60 is not). Used "this window" with no day-count instead — the exact span is already shown verbatim on screen (ReviewScreen's "Review · based on 1 Jul – 31 Aug 2026" line), so the narration doesn't need to also state a number here.
- **Fix 2's "nothing here is a habit yet" can't be used verbatim** — it contains the literal word "habit", which rule 8 (was rule 7 in v5) and `eval_narration.py`'s `check_state_fidelity` both specifically forbid for a no-habit narration (a hard-won rule from v3's failure, see section 3 above). Kept the requested *meaning* — this is a real, sufficiently-sized check that came up empty, not a shrug — without the banned word.
- **Fix 2's "across 101 trades" needs a number in the closing**, which v5's rule 4 banned outright, no exceptions. Rather than hardcode "101" (wrong for every other persona and impossible to keep in sync if the data ever changes) or drop the number entirely (losing the "this was a real check" weight the brief specifically wants), v6 carves out one narrow exception: the closing may

state summary.review_month.closed_trades verbatim, and only in the “nothing found” case. eval_narration.py’s check_number_tracing was updated to encode exactly this exception (a set-difference against {closed_trades} for the closing field only, gated on the same nothing_found condition the prompt uses) rather than loosened generally — every other field, and every other persona state, stays as strictly number-free as v5 enforced.

- **Fix 4b: checked every detector’s evidence for a fix-4b-shaped number**, not just size_up_after_loss.median_entry_size_otherwise is “your usual size” — the exact quantity size_up_after_loss’s rule refers to, so naming it is a direct substitution, not an inference. None of the other six rules refer to a quantity their own evidence contains: the fast-reentry window (5 min), the late-session cutoff (2:30pm) and the overtrading threshold (trade #7) are all detection parameters baked into detectors.py, not evidence fields — naming them in the rule text would mean inventing a number that isn’t in the analysis JSON, which the brief’s own constraint forbids. averaging_down and holds_losers_longer don’t have an evidence field that maps to their rule’s prescribed action (deciding size before entry; writing an exit level) the way “usual size” maps to size_up_after_loss’s rule. So only size_up_after_loss got the concrete-number treatment (narrate.py’s new RULE_BUILDERS); the other six keep their static RULE_TEMPLATES string, unchanged from v5. Verified against Arjun’s real evidence (median_entry_size_otherwise: 6578) before shipping: $\text{round}_{100}(6578) = 6600 \rightarrow$ “about ₹6,600”, matching the brief’s own example exactly.

6. Narration system prompt v6

(Note: an earlier, hypothetical “v6” was discussed and explicitly deferred in section 4 above — a stability-related prompt reinforcement that never shipped because the measured malformed-output rate didn’t warrant it. This is the first narration_system.md actually named v6; that discussion is unrelated to this one.)

Full prompt saved as prompts/narration_system.md (current). Changes from v5, all covered above: rule 4 gets one narrow numeric exception (nothing- found closing may state the trade count), a new rule 5 bans “this month” / “the month” for the evidence window in favour of “this window”, and the “closing depends on whether habits is empty” section is split into its two genuinely different empty-habits cases (insufficient-data vs. nothing- found), with the nothing-found case now specifying a two-sentence closing (explain what “nothing found” means at the trader’s scale, including the one permitted number $\rightarrow$ separate, quieter “check back” coda) instead of one.

Regeneration hit one real self-check failure, caught by the pipeline exactly as designed (not a prompt-wording failure — no v6.1 needed): arjun_revenge_sizer’s first two attempts both failed with check_number_tracing: orphan numbers: [6600.0]. Cause: narrate.py’s new_rule_size_up_after_loss (fix 4b, section 5 above) rounds median_entry_size_otherwise to the nearest ₹100 for the “about ₹6,600” phrasing — a deliberate, requested rounding — but eval_narration.py’s all_numbers() had no notion of “nearest-hundred rounding” as an allowed reformatting the way it already allowed a rate rounded to a percent. Fixed all_numbers() to also add the nearest-hundred form of any value ≥ 100 to its allowed set (mirrors the existing rate $\rightarrow$ percent special case exactly); re-ran and it passed on the first attempt, 0 retries. This is a case where the eval needed a fix for a legitimate new code-side number transformation, not the prompt or the generated prose — logged here since it happened during this prompt’s regeneration and blocked it until fixed.

Final regeneration, one Groq call per persona (all passed self-checks on the first attempt except the one arjun retry above), then eval_narration.py run for number tracing / advice filter / state fidelity as asked:

```

arjun_revenge_sizer: PASS
  number_tracing  [ok] ok
  advice_filter    [ok] ok
  state_fidelity   [ok] ok

neha_expiry_day: PASS
  number_tracing  [ok] ok
  advice_filter    [ok] ok
  state_fidelity   [ok] ok

sara_thin_data: PASS
  number_tracing  [ok] ok
  advice_filter    [ok] ok
  state_fidelity   [ok] ok

vikram_control: PASS
  number_tracing  [ok] ok
  advice_filter    [ok] ok
  state_fidelity   [ok] ok

ALL PASS

```

Stability was not re-run for v6 in the same session as the fixes above (only number tracing / advice filter / state fidelity were asked for then) — the pre-v6 stability run files were deleted rather than left in place, since comparing them to each other would validate v5’s wording, not v6’s, and a stale pass would have been more misleading than an honest “not run.”

Run separately, on request (python narrate.py --stability 5 --delay 20, 20 Groq calls): 19/20 succeeded first pass; neha_expiry_day run 3 hit the 8000 TPM free-tier limit after exhausting its own retries (a capacity question, not a content one) and was regenerated individually to complete the set. One self-check retry fired along the way — neha_expiry_day run 1 first came back with “a loss of ₹9,449” where the evidence says ₹449, the same digit-insertion flake recorded in section 4 above; the runtime self-check (not a human) caught it and regenerated automatically, so the saved run1 file is the corrected version, not the flaky one.

Result: **all four personas now have a full, genuine 5/5** — identical habit/watching ids and identical numbers across every rerun. This is better than any stability result this project has previously recorded: the v5-era numbers this section used to cite were 5/5 for arjun and neha but only a partial 3/5 (vikram) and 1/5 (sara), both cut short by a free-tier daily quota, not a content issue. eval_narration.py:

```

arjun_revenge_sizer: PASS
  number_tracing  [ok] ok
  advice_filter    [ok] ok
  state_fidelity   [ok] ok
  stability        [ok] 5 runs, identical ids and numbers

neha_expiry_day: PASS
  number_tracing  [ok] ok
  advice_filter    [ok] ok
  state_fidelity   [ok] ok
  stability        [ok] 5 runs, identical ids and numbers

sara_thin_data: PASS
  number_tracing  [ok] ok
  advice_filter    [ok] ok

```

```

state_fidelity    [ok] ok
stability         [ok] 5 runs, identical ids and numbers

vikram_control: PASS
number_tracing   [ok] ok
advice_filter     [ok] ok
state_fidelity    [ok] ok
stability         [ok] 5 runs, identical ids and numbers

ALL PASS

```

Vikram’s actual regenerated output (the one this whole prompt was about):

headline: “No repeated pattern met the cost or frequency thresholds in this window.” **closing:** “Across the 101 trades in this window, no single behaviour repeated enough or incurred enough cost to be confirmed, so the lack of a finding does not mean the period was successful. Check back at the next review to see if any new patterns emerge.”

Arjun’s regenerated rule (fix 4b), copied verbatim by the model from `fixed_rule` as instructed:

“After a loss, your next entry stays at your usual size — about ₹6,600. If you want to size up, it has to be a separate decision made before the session, not after a loss.”

7. Two copy fixes on the “nothing found” screen (Claude Code, 2026-09-12)

Two copy fixes on the “nothing found” screen. (1) Four ruled-out rows repeat the identical phrase “happened sometimes, but not consistently enough...”. Vary the phrasing so the actual reasons read differently (failed significance vs confounded vs too few). (2) In the closing line, “so the lack of a finding does not mean the period was successful” doesn’t follow logically from the first clause – split it into two sentences. Don’t change any numbers. Re-run `eval_narration.py` after.

Fix (1) was code, not the LLM — same as fix 1 in section 5: the $p=X \geq 0.01$ (could be noise) reason format was always mapped to one fixed sentence in `habitLabels.ts`’s `humanizeReason()`, regardless of how large X actually was. Vikram’s four “not statistically reliable” rows have $p = 0.764, 0.727, 0.263, 0.352$ — genuinely different degrees of “not a pattern” (0.764/0.727 are indistinguishable from pure chance; 0.263/0.352 are more borderline) — but all four got the identical phrase. Fixed by tiering the phrase on the p -value already in the string (5 tiers: $<0.05, <0.1, <0.3, <0.6, \geq 0.6$), still never displaying the number itself. This is not “make up 4 different reasons” — two of Vikram’s four (0.764, 0.727) are close enough that giving them the same tier, and therefore the same sentence, is the honest answer; forcing a 4th distinct phrase onto two p -values that are genuinely almost equal would have been the opposite mistake (fabricating a distinction the data doesn’t support). Result: 3 distinct phrases across Vikram’s 4 significance-based rows instead of 1.

Fix (2) is narration content (v6 → v7, `prompts/narration_system.md`). v6’s closing guidance asked for “no single behaviour repeated... **so** the lack of a finding does not mean the period was successful” as one acceptable shape — technically two ideas, but v6 never said they couldn’t be joined with “so”, and the model took that opening. The bug: “nothing was confirmed” doesn’t *cause* or *prove* “the window wasn’t fine” — they are two separate facts (a technical finding, and a caution against misreading it), not premise-and-conclusion. v7 makes this explicit: the closing is now three sentences, not two, with a standing instruction that sentences 1 and 2 must never be joined by “so”/“therefore”/“which means”/“meaning” because one doesn’t logically follow from the other. Regenerated (all four personas, since the prompt file is shared); Vikram’s new closing:

“Across the 101 trades in this window, no single behavior repeated often enough or incurred enough cost to be confirmed. That does not mean the window was successful; the loss shown elsewhere is not linked to a recurring, provable behavior. We’ll review the data again in the next cycle.”

No numbers were changed anywhere in either fix, per the brief — the trade count (101) is the same value from the same source (`summary.review_month.closed_trades`) as before; only the sentence structure and the ruled-out phrasing logic changed.

`eval_narration.py` (number tracing / advice filter / state fidelity, as asked — stale v6 stability run files deleted rather than left in place, same reasoning as section 6):

```
arjun_revenge_sizer: PASS
  number_tracing  [ok] ok
  advice_filter   [ok] ok
  state_fidelity  [ok] ok

neha_expiry_day: PASS
  number_tracing  [ok] ok
  advice_filter   [ok] ok
  state_fidelity  [ok] ok

sara_thin_data: PASS
  number_tracing  [ok] ok
  advice_filter   [ok] ok
  state_fidelity  [ok] ok

vikram_control: PASS
  number_tracing  [ok] ok
  advice_filter   [ok] ok
  state_fidelity  [ok] ok

ALL PASS
```

Fresh v7 stability check, run separately on request (`python narrate.py --stability 5 --delay 20, 20` Groq calls): 20/20 succeeded first pass — no rate limits, no self-check retries, cleaner than either prior stability run (v6’s had one rate-limit failure and one self-check retry; earlier v5-era runs never got all four to 5/5 at all). `eval_narration.py`:

```
arjun_revenge_sizer: PASS
  number_tracing  [ok] ok
  advice_filter   [ok] ok
  state_fidelity  [ok] ok
  stability        [ok] 5 runs, identical ids and numbers

neha_expiry_day: PASS
  number_tracing  [ok] ok
  advice_filter   [ok] ok
  state_fidelity  [ok] ok
  stability        [ok] 5 runs, identical ids and numbers

sara_thin_data: PASS
  number_tracing  [ok] ok
  advice_filter   [ok] ok
  state_fidelity  [ok] ok
  stability        [ok] 5 runs, identical ids and numbers
```

```
vikram_control: PASS
number_tracing [ok] ok
advice_filter [ok] ok
state_fidelity [ok] ok
stability [ok] 5 runs, identical ids and numbers

ALL PASS
```

v7 is now the second consecutive prompt version to reach a genuine 5/5 for all four personas (v6 did too, section 6) — the varied ruled-out phrasing (tiered by p-value, not free-form) and the three-sentence closing structure are both deterministic-enough templates that they didn't introduce new instability.

8. Live-app read-through catches a redundant-phrase bug (2026-09-12)

check the live app

(User pasted the actual rendered page text for all four personas from the live URL.) Reading the real output surfaced a bug none of the automated checks caught: Arjun's `why_it_matters` read "Those after-loss trades produced a loss of a loss of ₹1,95,375" — a duplicated phrase. Cause: `after_loss_net` is pre-formatted by `narrate.py`'s `_rupees_signed` into a complete phrase, "a loss of ₹1,95,375", not a bare number — the model wrote its own "produced a loss of" immediately in front of it anyway, producing the doubled phrase. `check_number_tracing` couldn't have caught this: the digit (195375) is correct and present either way; the defect is purely grammatical, not numeric.

Two other things checked and confirmed correct, not bugs, while reading the pasted text closely: - Sara's screen shows "Trades: 9" in the stat row but "18 closed trades" in the insufficient-data message — these are two different, intentional counts (`summary.review_month` = August only, 9 trades; the full 59-day evidence window = 18 trades), per the original brief's own design ("Review month... in three numbers... Evidence window stated explicitly"). Confirmed against `sara_thin_data_analysis.json` before concluding this, not assumed. - The pasted text shows "What we checked and ruled out (N)" with no bullets under it for any persona — expected, not a regression: `<details>` content isn't included in a plain select-all/copy of the rendered page unless it's expanded first, so this says nothing about whether fix 1 (section 5) is still rendering. Already verified separately via bundle inspection.

Fix, matching this project's own established pattern (a bug that slips past the eval suite gets a permanent new check, not just a one-off prompt patch — see the ₹9,449 digit-insertion incident, section 4): 1. `prompts/narration_system.md` v7 -> v8: rule 1 now explicitly names this failure mode ("Do not write 'produced a loss of a loss of ₹1,95,375' — a real failure this rule exists to prevent") and tells the model either to drop the pre-formatted value in as its own clause or introduce it with neutral wording ("at", "totalling", "amounting to") that doesn't repeat "loss"/"gain". 2. New `check_redundant_framing` in `eval_narration.py` — a regex for "a loss/gain of" immediately followed by another "a loss/gain of" — added as a 5th check to `eval_narration.py`'s per-persona report AND to `narrate.py`'s `SELF_CHECKS` (so a future occurrence is caught and retried automatically at generation time, not just flagged after the fact). Verified against the actual bug before regenerating: correctly fails on Arjun's pre-fix text (redundant framing found: `['a loss of a loss of']`) and passes clean on the other three (no false positives).

Regenerated all four personas against v8 (arjun and neha succeeded first call; vikram and sara hit the Groq TPM rate limit and were retried individually). Arjun's fixed text, first attempt, 0 self-check retries — the prompt guardrail worked without needing the runtime retry:

“The cost of this behaviour is ₹1,05,732. It occurred in 96 trades, with a median entry size after a loss of ₹12,562 versus ₹6,578 otherwise, giving a size ratio of 1.91x. The after-loss win rate was 28% compared with 39% otherwise, and those trades were a loss of ₹1,95,375.”

eval_narration.py (now 4 non-stability checks per persona):

```
arjun_revenge_sizer: PASS
  number_tracing      [ok] ok
  advice_filter        [ok] ok
  state_fidelity       [ok] ok
  redundant_framing    [ok] ok

neha_expiry_day: PASS
  number_tracing      [ok] ok
  advice_filter        [ok] ok
  state_fidelity       [ok] ok
  redundant_framing    [ok] ok

sara_thin_data: PASS
  number_tracing      [ok] ok
  advice_filter        [ok] ok
  state_fidelity       [ok] ok
  redundant_framing    [ok] ok

vikram_control: PASS
  number_tracing      [ok] ok
  advice_filter        [ok] ok
  state_fidelity       [ok] ok
  redundant_framing    [ok] ok

ALL PASS
```

Stale v7 stability run files deleted (same reasoning as sections 6-7 — they validate v7’s wording, not v8’s). Not re-run in this pass; ask for a fresh one the same way as before if wanted.

9. v8 stability check: daily quota blocks a full run (2026-09-12)

Run the 5x stability check against v8 for all four personas and give me the results.

Attempted `python narrate.py --stability 5 --delay 20` (20 Groq calls). Got through arjun’s full 5 and one of neha’s before hitting Groq’s *daily* token quota (200,000 TPD) — a harder wall than the per-minute limit hit earlier in this project, not something a short retry clears. The day’s earlier v6 and v7 stability batches plus several individual regenerations had already used most of the day’s budget before this attempt started.

What completed, exactly: - **arjun_revenge_sizer: 5/5**, identical habit ids and identical numbers across all five runs. - **neha_expiry_day: 1/5**. That single run’s first attempt tripped the new `check_redundant_framing` self-check from section 8 — Groq regenerated “a loss of a loss of” again, on a live call, and `narrate_persona`’s runtime retry caught it and produced a corrected version before the run was ever saved. This is direct evidence the v8 guardrail works against a real recurrence, not just the original instance — but one run is not a stability result (nothing to compare it against), so it isn’t counted as a stability pass. - **vikram_control: 0/5. sara_thin_data: 0/5**. The daily quota was exhausted before either got a single successful call in this batch.

Option 3 [accept partial evidence for now]. Update EVALS.md to state the v8 stability position honestly: Arjun 5/5 identical ids and numbers; Neha 1/5 completed, with that run tripping the redundant-framing self-check and regenerating automatically before save; Vikram and Sara not re-run against v8 due to Groq daily quota, with their earlier v6 runs and the structural argument (no numbers in prose) noted. Make sure eval_narration.py reports "no stability runs found" as SKIP, not PASS. Add the "a loss of a loss" bug to the iteration log, including that automated checks missed it and a human read caught it. Commit and push. Then stop – no more test runs tonight.

A real bug found while doing this, separate from the narration content itself: eval_narration.py's main() coerced check_stability's None ("no stability runs found") into True before storing it in the results dict, so the per-check print loop's own "skip" if ok is None else ... branch was dead code for stability specifically — a persona with zero stability runs printed stability [ok] no stability runs found, visually indistinguishable from a genuine pass unless you read the message text closely. Fixed by storing the raw tri-state result instead of coercing it; Vikram and Sara now correctly print [skip]. This never affected overall PASS/FAIL (the persona_ok calculation already excluded None from consideration and separately checked for a real False), only what was displayed — but "displayed as ok" and "is ok" need to actually agree, and they didn't.

eval_report.md and README.md updated to state this exact position — Arjun's real 5/5, Neha's single data point (not a stability claim), Vikram/Sara's honest gap plus the v6 result and structural argument that still stand independently of it — rather than a summary number that overstates what actually ran tonight.

No further Groq calls made after this, per the instruction to stop for the night.

10. ONE_PAGER.md and the submission PDF (2026-09-12)

Add a new section to ONE_PAGER.md called "What I'd validate first", placed after "What testing changed" and before "Scope". Keep the whole one-pager under 700 words — cut from other sections if needed.

The section must state plainly that this product rests on assumptions I have not tested with real traders, then list the three riskiest ones, each with how I'd test it and what result would kill or change the product:

1. Desirability — that a trader who loses money wants a diagnosis rather than avoiding the review entirely. Test: does the review get opened after a losing month, or only after a winning one?
2. Behaviour change — that naming one habit and getting a commitment actually reduces that habit next month. Test: habit recurrence rate for committed vs non-committed users.
3. Trust — that a trader accepts a behavioural verdict from his own broker, given the broker earns on his volume.

Then add a short "How I'd measure success" paragraph: the metric is habit recurrence next month, not opens or session time; the counter-metric is churn, because if diagnosed traders leave rather than improve the thesis is wrong; and note the noise problem, that a trader can fix a real habit and still have a worse month.

Write it in the same plain voice as the rest of the document. Do not invent user quotes, interviews, or research findings. Commit and push, then rebuild TradeLens_Submission.pdf.

ONE_PAGER.md and TradeLens_Submission.pdf didn't exist anywhere in this repo, its git history, or elsewhere on the filesystem checked before this prompt — asked rather than guessed at content or a build pipeline that might exist somewhere out of view. The user pasted the exact ONE_PAGER.md content to save verbatim (744 words) and asked for a new pandoc-based PDF pipeline; that baseline was committed first, separately, per the user's own instruction, before this section was added.

The 700-word budget was tight from the start: the baseline alone was 744 words, over budget before adding anything. The two new sections, written as economically as possible while still covering every required element (the assumptions statement, three risks each with a test and a kill/change signal,

and the metric/counter-metric/noise paragraph), needed real cuts elsewhere — not light copy-editing but removing whole clauses and one full sentence (the “duplicated phrase” anecdote in “What testing changed”, now that “What I’d validate first” carries a similar epistemic-honesty note). Two supporting SEBI statistics (the 92%-of-losses and 59%-same-day-expiry figures) were also cut from “The problem” as the least load-bearing facts, keeping the core argument (87.7% lose, 90% repeat) intact. Final: 699 words by `wc -w` — confirmed directly against the assembled file, not estimated from section subtotals (which undercounted by not including the “##” heading tokens themselves).

No user quotes, interviews, or research findings were invented — the three “Test:” lines are proposed methods for validating the stated assumptions, matching the section’s own framing (“what I’d validate first”), not claimed results.

PDF pipeline: pandoc (installed via winget — a LaTeX distribution was too heavy, and wkhtmltopdf’s installer needed UAC elevation this sandboxed shell couldn’t grant) converts the markdown to an HTML fragment; a new `scripts/build_pdf.js` wraps it in print-styled HTML (plain typography, the app’s own amber accent, matching the product’s restrained design voice) and Puppeteer (a plain npm devDependency, no elevation needed, ships its own Chromium) prints it to `TradeLens_Submission.pdf`. `npm run build:pdf` from the repo root. Verified by reading the actual rendered PDF, not just checking the file was written: two pages, correct section order, no rendering artifacts.

ARTIFACT 3 OF 4

The evals

TradeLens — Detector Evals

The question

Would a trader trust this with their own money story? That breaks into two failures to avoid:

1. **Missing a real habit.** The review is useless.
2. **Inventing a habit that isn't there.** The review is actively harmful, and trust ends the first time it happens.

Failure 2 is worse, so the thresholds are set to protect against it first.

Method

- **Ground truth.** Four synthetic traders with habits planted at known strength, including a control with none and a thin-data trader (see `habit_library.md`).
- **No tuning on the test set.** Thresholds were tuned on dev seeds 1–30. The results below use held-out seeds 101–130, which were never looked at during tuning.
- **Pass criteria per persona-month.** Planted personas: every planted habit found and nothing unplanted reported. Control: zero habits. Thin data: the “not enough data” response.

Held-out results (30 simulated traders per persona)

Persona	1-month window	60-day window
Arjun: both planted habits found, nothing invented	13/30	28/30
Neha: expiry-day as #1, nothing invented	20/30	29/30
Vikram (control): zero habits	30/30	30/30
Sara (thin data): says “not enough data”	30/30	30/30
Wrong habit reported as #1 (worst failure)	2/30 (Neha)	0/30

Almost every 1-month failure is a **miss**, not a false alarm, and most misses surface as *watching* rather than disappearing.

Why 60 days: detection power (dev seeds)

Habit (planted effect)	Detected, 1 month	Detected, 2 months
Size-up after loss (2× size)	28/30	30/30
Averaging down (added lots −12%)	14/30	27/30
Expiry-day, −10% per trade	1/30	5/30
Expiry-day, −18% per trade	7/30	16/30
Expiry-day, −28% per trade	18/30	26/30
Expiry-day, −39% per trade	26/30	30/30
Control, false alarms	1/30	0/30

Lottery-like payoffs (cheap expiry-day options) are high-variance. One month of trades often can't separate a costly habit from a bad month. Loosening the threshold would buy recall with false alarms (see the sweep below). A longer evidence window buys recall for free.

Threshold sweep, 1-month, dev seeds:

p-value threshold	Arjun	Neha	Control correct
0.01 (<i>chosen</i>)	12/30	18/30	29/30
0.02	15/30	19/30	29/30
0.05	21/30	21/30	28/30
0.10	26/30	23/30	27/30

What broke, and what changed (iteration log)

1. **Missed Neha's planted habit ($p = 0.46$).** Significance was tested on ₹ P&L. Her expiry-day trades were small in ₹, so noise from her bigger normal trades drowned the signal. *Fix*: test on size-normalised return and report cost in ₹.
2. **Averaging down vanished for Arjun.** A rule merged habits that shared trades, and since about half his trades follow a loss, averaging down kept getting absorbed into size-up. *Fix*: distinct behaviours are never merged; only time/day patterns are tested for confounding.
3. **Neha got "overtrading" as #1 in some months.** Her 7th-plus trades are all on expiry days, so overtrading was a proxy. *Fix*: a two-way test. Overtrading doesn't hold once expiry trades are removed; expiry-day trading does hold once 7th-plus trades are removed.
4. **Control false positive worth ₹1,145.** A flat ₹1,000 bar is meaningless for someone buying ~₹15 lakh of premium a month. *Fix*: the bar became the larger of ₹1,000 and 1% of premium bought.
5. **Implausible data.** Arjun lost ₹3.7 lakh in a month, and his size-up drill-down showed averaging-down trades, which muddled two stories. *Fix*: realistic sizing, and examples exclude trades claimed by another habit.
6. **The "second month" was September 2026, which is in the future.** *Fix*: the window became July + August.

Product decisions that came from the evals

- **60-day evidence window**, with the review still framed monthly.
- **"Watching" state** for material but unproven patterns, so the product can be honest without being silent.
- **"Holds losers longer" stays uncoded.** Order history can't prove its cost, so it's never shown with a ₹ figure.
- **Habit costs are never added together**, because the same trade can appear in two habits.

Ship bar (detector layer)

Measured on held-out seeds, 60-day window: - planted habit recovered in $\geq 90\%$ of months; - control false alarms $\leq 5\%$; - thin data handled correctly 100% of the time; - zero wrong-habit calls.

Met at 60 days (93–97%, 0%, 100%, 0). **Not met** at one month, which is why the window changed.

Ship bar (narration layer)

Zero orphan numbers, zero advice-word hits, zero habit-language misuse, identical output across reruns — across all four personas.

Met on number tracing / advice filter / state fidelity / redundant framing: 4/4 personas, every generation, on the current prompt (v8 — see PROMPTS_LOG.md sections 6-8 for the full version history: window wording and the concrete `size_up_after_loss` rule in v6; varied ruled-out phrasing and a “nothing found” closing that doesn’t chain two unrelated ideas with “so” in v7; a redundant-phrase guardrail, found by reading the live app rather than by any automated check, in v8).

Stability against v8, honestly: incomplete, not a full 5/5. A fresh 5x-per-persona batch (20 Groq calls) was started and ran into Groq’s *daily* token quota (200,000 TPD, not the per-minute limit) partway through, on top of the day’s earlier v6/v7 stability batches and several individual regenerations. What actually completed:

- **Arjun: a genuine 5/5** — 5 runs completed, identical habit ids and identical numbers across all of them.
- **Neha: 1/5 completed**, not 5. That single run is worth noting on its own terms, not just as a partial count: it initially tripped the new `check_redundant_framing` self-check (the same “a loss of a loss of ₹X” failure mode from section 8) and `narrate_persona`’s runtime retry caught it and regenerated automatically before the run was saved — the guardrail added in v8 working against a live recurrence, not just the original bug. One run cannot establish stability by itself (there’s nothing to compare it against), so this is not treated as a stability pass, only as a data point.
- **Vikram and Sara: not re-run against v8 at all** — the daily quota was exhausted before either got a single call in this batch. Their last *verified* stability result is the v6 run (section 6): a genuine 5/5 for both. That result doesn’t speak to v8’s wording, but the structural argument from section 6 still holds independently of any specific prompt version: v8, like every version since v5, makes `headline` and `closing` unconditionally number-free for both of them (Vikram’s one narrow trade-count exception is the sole deliberate exception, itself copied verbatim from a fixed JSON field, not authored by the model), and forces `habits/watching/also_noticed` to all be empty whenever `insufficient_data` is true or nothing qualified — which is Vikram and Sara’s situation every time. There is very little surface for a prompt wording change to introduce fresh instability into a screen with no numbers and no habit content to vary.

Re-run `python narrate.py --stability 5 --delay 20` then `python eval_narration.py` once Groq’s daily quota resets to get a real 5/5 for all four against v8 — this section will say so plainly when that happens, not before.

Narration layer evals

Provider/model: Groq, openai/gpt-oss-120b, temperature 0.2 (see README for why Groq — free tier, no card — and how to swap providers). Five automated checks in `eval_narration.py`, run against all four personas:

Check	What it verifies	Result
Number tracing	every ₹/%/ratio/count in the narration traces to a value in the analysis JSON	PASS , all 4
Advice filter	no instrument/strike/direction/prediction words anywhere	PASS , all 4

Check	What it verifies	Result
State fidelity	no habit-language for Vikram/Sara; “watching” never phrased as confirmed	PASS , all 4
Redundant framing (v8)	no “a loss/gain of” immediately followed by another “a loss/gain of” (a real generation duplicated one — see PROMPTS_LOG.md section 8)	PASS , all 4
Stability (5 reruns)	identical habit/watching ids and identical numbers across 5 regenerations	Arjun: PASS, genuine 5/5. Neha: 1/5 completed (eval_narration.py reports [ok] 1 runs — technically true of the one run it has, not a real stability claim). Vikram/Sara: SKIP — 0 runs against v8, correctly reported as [skip] no stability runs found, not as a pass (see the fix below)

eval_narration.py used to coerce “no stability runs found” into a printed [ok], indistinguishable from a genuine pass — fixed so it now reports [skip] honestly (Vikram and Sara currently show this). This only changes what’s *displayed*; a persona was never marked overall PASS or FAIL based on a missing stability run either before or after the fix, only on whether a run that did happen showed instability.

Every single-shot generation, on every prompt version from v2 onward, has independently passed number-tracing and state-fidelity regardless of whether a full 5x stability batch was current at the time — the thing stability is checking for (the model violating a numeric or empty-habits rule) is already caught per-run by the checks above, and by narrate_persona’s own self-check retry at generation time. Remember that a stability run in results/narration_stability/ only speaks to the prompt it was generated against — re-run it after any prompt edit rather than trust a stale pass; see PROMPTS_LOG.md section 9 for the full v8 stability attempt and what blocked it.

Not yet tested

- **Detectors without a dedicated persona.** Late-session, fast re-entry (only tested as a confounder) and holds-losers-longer.
- **Real-world validity.** Synthetic data proves the system recovers what was planted, not how common these habits are among Nubra users. Planted effect sizes are assumptions.

How to run

```
python run_evals.py heldout 2      # ship-decision table (60-day window)
python power.py                   # detection power (dev seeds)
python tune.py                    # threshold sweep (dev seeds)
python export_demo.py              # demo order histories + analysis JSON for the app
```

ARTIFACT 4 OF 4

The working app

Live prototype: <https://app-azure-zeta-42.vercel.app> — no login required.

Source code: <https://github.com/Shraman123/tradelens>

TradeLens

DEMO PERSONA

Arjun

Sizes up after a loss, and averages down on losing positions.

Confirmed habit + a second pattern still only "watching"

Neha

Cheap expiry-day options are where most of the damage...

One confirmed habit

← Back to review

HABIT #1

Sizes up on the trade after a loss

strong confidence

WHAT THIS COST YOU

₹1,05,732

Net P&L if after-loss trades had used your usual entry size

EVIDENCE

After-loss trades	96
Median entry size after a loss	₹12,562
Median entry size otherwise	₹6,578
Size ratio	1.91x
Net P&L on after-loss trades	-₹1,95,375
Win rate after a loss	28%
Win rate otherwise	39%

RULE FOR NEXT MONTH

After a loss, your next entry stays at your usual size — about ₹6,600. If you want to size up, it has to be a separate decision made before the session, not after a loss.

EXAMPLE TRADES

NIFTY 21JUL26 24200 PE

16 Jul 2026 · 11:25

-₹8,934

View raw orders

NIFTY 04AUG26 24950 CE

29 Jul 2026 · 11:51

-₹7,160

View raw orders

NIFTY 11AUG26 24900 CE

10 Aug 2026 · 12:38

-₹6,569

View raw orders

NIFTY 25AUG26 25000 PE

19 Aug 2026 · 11:08

-₹5,999

View raw orders

NIFTY 21JUL26 24200 PE

20 Jul 2026 · 11:20

-₹5,789

View raw orders

Habit costs overlap and must never be added together.

Arjun -- Habit Detail (confirmed habit, evidence, and the concrete rule)

27 / 30

DEMO PERSONA

Arjun

Sizes up after a loss, and averages down on losing positions.

Confirmed habit + a second pattern still only "watching"

Neha

Cheap expiry-day options are where most of the damage...

One confirmed habit

Review · based on 1 Jul – 31 Aug 2026

Expiry-day trades are costing you a lot.

NET P&L

-₹94,438

TRADES (AUG)

99

WIN RATE

31%

HABIT #1

strong confidence

Losses on expiry-day (ODTE) trades

₹1,16,891

Expiry-day trades cost you ₹1,16,891 across 90 trades, resulting in a loss of ₹1,16,891 and an average loss per trade of ₹1,299, with a win rate of 29% and an average return of a 28% loss. The remaining 124 trades showed a loss of ₹449 on average, a 5% loss overall, and a win rate of 37%.

RULE FOR NEXT MONTH

Set a hard cap on expiry-day trades before the open.

[View evidence & example trades >](#)

What we checked and ruled out (6)



Consider applying the rule to limit expiry-day trades in the next review period.

[Commit to a rule for next month](#)

Neha -- Review (one confirmed habit, ranked and costed)

Neha

Cheap expiry-day options are where most of the damage...

One confirmed habit

Vikram

Losing money, but no repeatable pattern in how he trades.

Nothing found — the product says so, honestly

Review · based on 1 Jul – 31 Aug 2026

No repeated, cost-significant pattern emerged in this window.

NET P&L

-₹52,169

TRADES (AUG)

101

WIN RATE

43%

What we checked and ruled out (7)

Sizes up on the trade right after a loss — the size difference was within normal variation

Adds to positions that are already losing — didn't happen this window

Jumps back in within minutes of a loss — happened about as often as plain chance alone would predict — not a pattern

Expiry-day trades are where the money goes — happened about as often as plain chance alone would predict — not a pattern

Late-session trades lose money — showed up more than a coincidence alone might suggest, but not consistently enough to call it a pattern

Trades after the 6th of the day lose money — showed up sometimes, but not consistently enough across your trades to call it a pattern

Holds losers much longer than winners — how long you held winners versus losers wasn't meaningfully different

Across the 101 trades this window, no single behaviour repeated often enough or cost enough to be proven from the order history. That does not mean the window was successful; the loss shown elsewhere is not linked to any recurring, provable behaviour. Check back at the next review for new insights.

Vikram -- Review, nothing found ("What we checked and ruled out" expanded)

DEMO PERSONA

Vikram

Losing money, but no repeatable pattern in how he trades.

Nothing found — the product says so, honestly

Sara

Only a handful of trades so far this window.

Not enough data to call anything a habit

Review · based on 1 Jul – 28 Aug 2026

Not enough trades to identify a repeatable behaviour

18 closed trades; need at least 20 before calling anything a habit.

NET P&L

-₹12,352

TRADES (AUG)

9

WIN RATE

11%

We don't have enough trading history yet to spot a consistent behaviour. Keep building your trade record, and we'll revisit this in the next review.

Sara -- Review, not enough data